

# Reproducing Slow-leader and Slow-follower: One-Size-Fits-None (NSDI'25)

## Slow Follower & Slow Leader — Simple Step-by-Step

### 1) Requirements & Install (once)

# Docker + compose plugin (Ubuntu)

```
sudo apt-get update
```

```
sudo apt-get install -y ca-certificates curl gnupg lsb-release
```

```
sudo install -m 0755 -d /etc/apt/keyrings
```

```
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo gpg --dearmor -o  
/etc/apt/keyrings/docker.gpg
```

```
echo "deb [arch=$(dpkg --print-architecture) signed-by=/etc/apt/keyrings/docker.gpg]  
https://download.docker.com/linux/ubuntu $(lsb_release -cs) stable" \
```

```
| sudo tee /etc/apt/sources.list.d/docker.list > /dev/null
```

```
sudo apt-get update
```

```
sudo apt-get install -y docker-ce docker-ce-cli containerd.io docker-buildx-plugin  
docker-compose-plugin
```

# Optional helpers

```
sudo apt-get install -y tmux jq coreutils
```

### 2) Bring up the cluster

# From folder with your docker-compose for Cassandra/etcd

```
docker compose up -d
```

# Check containers are running

```
docker ps
```

### 3) Slow Follower Experiment

In this fault, one follower node is made slow via tc netem (delay or loss) while the leader is healthy.

We use `run_netdelay.sh` (mid-run injection) or `run_netfault.sh` (full-run injection).

#### (a) Baseline

```
./run_netfault.sh loss "" "" 60 baseline
```

No fault; records baseline throughput.

#### (b) Delay — Full Run

```
bash
```

```
CopyEdit
```

```
./run_netfault.sh delay "cassandra_b" "100ms" 60 delay100_full
```

(c) Delay — Mid Run

bash

CopyEdit

```
./run_netdelay.sh "cassandra_b" "100ms" 60 delay100_mid 30
```

(d) Loss — Full Run

bash

CopyEdit

```
./run_netfault.sh loss "cassandra_b" "20%" 60 loss20_full
```

(e) Loss — Mid Run

bash

CopyEdit

```
./run_netdelay.sh "cassandra_b" "20%" 60 loss20_mid 30
```

#### 4) Slow Leader Experiment

Here, the leader node is slowed to test how leadership bottlenecks affect throughput.  
The target leader container must be identified first.

(a) Find current leader

bash

CopyEdit

```
docker exec etcd0 etcdctl endpoint status -w table \
--endpoints="http://etcd0:2379,http://etcd1:2379,http://etcd2:2379"
```

Look for Leader column; note the leader's container name.

(b) Delay — Full Run

bash

CopyEdit

```
./run_netfault.sh delay "<leader_container>" "100ms" 60 leader_delay100_full
```

(c) Delay — Mid Run

bash

CopyEdit

```
./run_netdelay.sh "<leader_container>" "100ms" 60 leader_delay100_mid 30
```

(d) Loss — Full Run

bash

CopyEdit

```
./run_netfault.sh loss "<leader_container>" "20%" 60 leader_loss20_full
```

(e) Loss — Mid Run

bash

CopyEdit

```
./run_netdelay.sh "<leader_container>" "20%" 60 leader_loss20_mid 30
```

#### 5) Sweep runs (optional)

You can automate multiple runs for a range of values.

Loss sweep example:

```
bash
CopyEdit
./sweep_loss.sh 60
```

This runs baseline + loss {1%,5%,...,70%} against cassandra\_b & cassandra\_c.

#### 6) Outputs per run

- Cassandra stress log → outputs/<timestamp>\_<label>/<label>\_cassandra-stress.log
- metadata.txt → fault parameters + run timestamp.
- You can use tail\_all.sh to quickly preview last 15 lines of all logs.

```
bash
CopyEdit
./tail_all.sh
```

#### 7) Reset / clear injection

```
bash
CopyEdit
docker exec <container> tc qdisc del dev eth0 root 2>/dev/null || true
```

#### 8) Copy results to local machine

```
powershell
CopyEdit
scp -i "C:\path\to\key.pem" -r ^
cc@<HOST>:~/outputs ^
C:\path\to\local\outputs
```

#### 9) Collect your result folders

Decide where your outputs/<timestamp>\_<label>/ directories live (on the remote or copied locally).

Examples:

- Linux: ~/cassandra-demo/outputs
- Windows (after scp): C:\Newer-fault\outputs

In the plotting code below, you will point BASE\_DIR to this folder.

## 10) Parse & plot throughput from cassandra-stress logs

Use this Python script to extract ops/s over time from each run's log and produce line charts. It assumes each run has:

php-template

CopyEdit

outputs/<stamp>\_<label>/<label>\_cassandra-stress.log

The regex targets typical cassandra-stress periodic lines reporting throughput. If your image prints a slightly different format, adjust the regex noted in the comments.

```
# save as plot_throughput.py and run: python3 plot_throughput.py
```

```
# Requirements: Python 3.10+, pandas, matplotlib
```

```
import os, re, sys
```

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
# 1) POINT THIS TO YOUR outputs/ DIRECTORY
```

```
# Example Linux: BASE_DIR = "/home/cc/cassandra-demo/outputs"
```

```
# Example Windows: BASE_DIR = r"C:\Newer-fault\outputs"
```

```
BASE_DIR = "/home/cc/cassandra-demo/outputs"
```

```
# 2) A helper to parse a single stress log and return DataFrame(second, ops)
```

```
# Adjust the regex if your logs differ. We look for "...op/s" numbers.
```

```
LINE_RX = re.compile(r'(?i)\b([0-9]+(?:\.[0-9]+)?)\s*op/s\b')
```

```
def parse_stress_log(path):
```

```
    ops = []
```

```
    sec = 0
```

```
    with open(path, 'r', encoding='utf-8', errors='ignore') as f:
```

```
        for line in f:
```

```
            m = LINE_RX.search(line)
```

```
            if m:
```

```
                try:
```

```

        val = float(m.group(1))

    except:

        continue

    ops.append((sec, val))

    sec += 1

if not ops:

    return pd.DataFrame(columns=["second", "ops"])

df = pd.DataFrame(ops, columns=["second", "ops"])

return df

# 3) Collect all runs

runs = [] # list of (label, df)

for d in sorted(os.listdir(BASE_DIR)):

    run_dir = os.path.join(BASE_DIR, d)

    if not os.path.isdir(run_dir):

        continue

    # label is the suffix after last underscore

    label = d.split("_", maxsplit=1)[-1] if "_" in d else d

    logfile = os.path.join(run_dir, f"{label}_cassandra-stress.log")

    if not os.path.isfile(logfile):

        continue

    df = parse_stress_log(logfile)

    if not df.empty:

        # light smoothing for readability (moving average 5)

        df["ops_smooth"] = df["ops"].rolling(window=5, min_periods=1,
        center=True).mean()

        runs.append((label, df))

if not runs:

```

```

print("No parsable runs found under:", BASE_DIR)

sys.exit(0)

# 4) Plot each run individually to PNG for slides

PLOTS_DIR = os.path.join(BASE_DIR, "_plots")

os.makedirs(PLOTS_DIR, exist_ok=True)

for label, df in runs:

    plt.figure(figsize=(10,4))

    plt.plot(df["second"], df["ops_smooth"], label=label)

    plt.title(f"Throughput vs Time — {label}")

    plt.xlabel("Time (s)")

    plt.ylabel("Ops/s")

    plt.grid(linestyle=":")

    plt.legend()

    out = os.path.join(PLOTS_DIR, f"{label}_throughput.png")

    plt.tight_layout()

    plt.savefig(out, dpi=150)

    plt.close()

    print("Saved:", out)

# 5) (Optional) Quick overlay helpers by family (Full/Mid, Loss/Delay)

def overlay(pattern, title, legend_title):

    import fnmatch

    subset = [(lbl, df) for lbl, df in runs if fnmatch.fnmatch(lbl, pattern)]

    if not subset:

        print("No matches for pattern:", pattern); return

    plt.figure(figsize=(12,5))

    for lbl, df in subset:

```

```

plt.plot(df["second"], df["ops_smooth"], label=lbl)

plt.title(title); plt.xlabel("Time (s)"); plt.ylabel("Ops/s")

plt.grid(linestyle=":")

plt.legend(title=legend_title, loc="center left", bbox_to_anchor=(1,0.5))

out = os.path.join(PLOTS_DIR, f"{title.replace(' ', '_')}.png")

plt.tight_layout(); plt.savefig(out, dpi=150); plt.close()

print("Saved:", out)

# Examples (tweak to your labels):

overlay("*loss*full", "Loss Full — Throughput Overlays", "Runs")

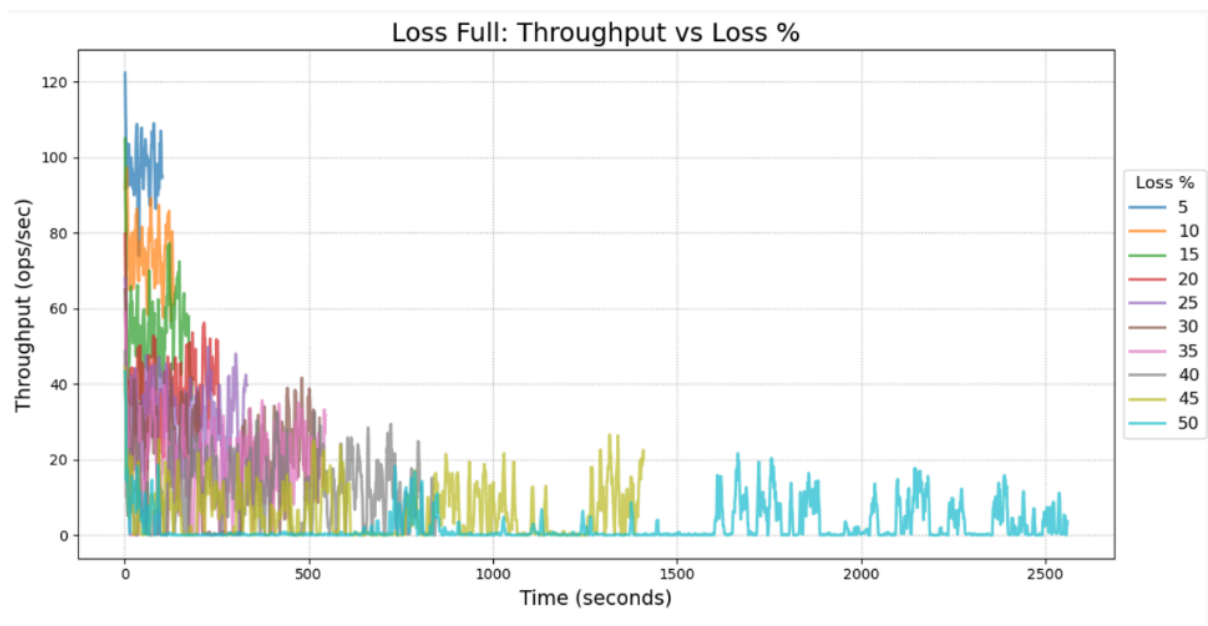
overlay("*loss*mid", "Loss Mid — Throughput Overlays", "Runs")



overlay("*delay*full", "Delay Full — Throughput Overlays", "Runs")

overlay("*delay*mid", "Delay Mid — Throughput Overlays", "Runs")


print("\nAll plots stored under:", PLOTS_DIR)

```









## 10) Run Example

`./run_slowfollower.sh 10% 120 loss10_mid loss 60`

```
(blockade-venv) cc@osfn-cc:~/cassandra-demo$ ./run_slowfollower.sh 10% 120 loss10_mid loss 60
[INFO] Label : loss10_mid
[INFO] Fault Mode/Value : loss 10%
[INFO] Duration : 120s
[INFO] Inject at : 60s
[INFO] Target follower : etcd1
[INFO] Output dir : etcd_bench_results/20250727_200645_loss10_mid
[INFO] Launching 10000 puts (10 parallel) via etcd1 ...
[INFO] Starting throughput monitor...
[INFO] Injecting loss 10% on etcd1 at 60s...
[INFO] Cleaning up NetEm on etcd1...
Latency median (p50) : 107 ms
Latency 95th percentile : 222 ms
Latency 99th percentile : 327 ms
Latency max : 1549 ms
Total puts : 10000
Elapsed (ms) : 116817
Avg throughput (ops/sec) : 85
[SUCCESS] Results in etcd_bench_results/20250727_200645_loss10_mid
```

`./run_slowfollower.sh 50ms 60 slowfollower_mid delay 30`

```
(blockade-venv) cc@ostn-cc:~/cassandra-demo$ ./run_slowfollower.sh 50ms 60 slowfollower_mid delay 30
[INFO] Label : slowfollower_mid
[INFO] Fault Mode/Value : delay 50ms
[INFO] Duration : 60s
[INFO] Inject at : 30s
[INFO] Target follower : etcd1
[INFO] Output dir : etcd_bench_results/20250727_190940_slowfollower_mid
[INFO] Launching 10000 puts (10 parallel) via etcd1 ...
[INFO] Sleeping 30s before injecting fault into etcd1 ...
[INFO] Injecting delay=50ms on etcd1 ...
[INFO] Removing NetEm from etcd1 ...
Latency median (p50) : 158 ms
Latency 95th percentile : 179 ms
Latency 99th percentile : 188 ms
Latency max : 198 ms
Total puts : 10000
Elapsed (ms) : 148024
Throughput (ops/sec) : 67
[SUCCESS] Slow-follower run complete. Results in etcd_bench_results/20250727_190940_slowfollower_mid
```